

Report iteration 3 - SWOP

Dieter Geboers, Wouter Schaekers, Stefaan Truijen

-

Professor: Tom Holvoet

Project Advisor: Marko van Dooren

2012-04-06

1 Preface

This iteration is the final one. The new assignment instructed us to update our project so that it is able to support multiple campuses on hospitals. With this new feature came a new feature for doctors as well. Doctors gained the ability to travel between campuses. They also get a say in their behaviour concerning travelling between campuses: go back and forth between all campuses, or rather have an own schedule concerning locations?

Preemptive scheduling and tasks that support multiple priority levels were not part of our assignment as we're a team of three members.

We delivered a report on the refactoring changes we planned to execute during this iteration as well as two test scenarios for our system. However, as we changed the way the system works, the test scenarios were forced to be adapted. We also thought that it would be best to create an additional testing scenario to prove that the "basics" are not the only thing we can do.

Our project advisor, Willem Penninckx, was replaced by Marko van Dooren for this iteration due to illness. We had some very constructive meetings with them both and are very thankful for the insight they gave us. Especially when we started to deviate from the core of the project by over thinking some of the design decisions.

For the first time in our SWOP history we can say with confidence that everything works like it should and there are no missing features.

Contents

1	Preface	1
2	Introduction	4
3	Package visible constructors	4
4	Creational design patterns	5
4.1	User factories and user manager	5
4.2	Machine builder and machine pool	6
4.3	Medical test factories and patient files	7
4.4	Treatment factories and diagnosis	7
4.5	Diagnosis and patient files	7
4.6	Warehouse builder and campus	8
5	Task	8
5.1	Generic tasks versus a hierarchy of tasks	8
5.2	The state of a Task	10
5.3	The description of a Task	10
6	Warehouse	10
7	Scheduling	11
7.1	Description and Requirements	11
7.2	Scheduler and TaskManager	12
7.3	Schedulable and Requirable	13
7.4	HospitalDate	13
7.5	TimeTable	13
7.6	TimeSlot	14
7.7	TimePoint	14
7.8	Impact of having different locations	14
7.9	The location preference for doctors	14
7.9.1	getLocationAt()-method	15
8	The updated testing scenario	15
8.1	The state of the hospital at initialisation	15
8.2	The events that occur	16
9	Things that can be done better	16
9.1	Diagnose state	16
9.2	Comments	16
9.2.1	TimeTables & Locations	16
9.3	Factories & consistency	17
9.4	DoctorState	17
9.5	Results	17
10	Eclemma and unit tests	17

11 Capsules	17
12 Conclusion	17

2 Introduction

During the feedback session of iteration two, some major (as well as minor) issues with our code and design floated to the surface:

- MachinePool had hard coded methods to create different types machines;
- There were no system sequence diagrams available;
- The UML of the domain layer did not have any relevant methods on it;
- The implementation of Warehouse was lacking in a lot of aspects:
- When an activity that required something from a warehouse got scheduled, there were no items removed from that warehouse;
- The way Warehouse kept track of its items and their maximum capacity was hardcoded;
- The system for creating stock orders and delivering them to the warehouse wasn't implemented successfully;
- ...
- The unit for time (milliseconds) used in the system was nowhere documented.
- Some methods and controllers required the same set of parameters and instead of an object that groups them together, they were all mentioned explicitly.
- ...

We have taken all of the above (and more) into account while implementing iteration three to make certain that our hard work would not have been in vain. In this report we'll motivate the necessity of changes we have put the project through. Also, we'll try to give some other solutions for the problems we encountered that caused us to create the system as it is. We'll then draw a conclusion and leave it at that.

As for the amount of time we spent working on this iteration in detail, we refer to <http://code.google.com/p/swop-dswx/wiki/WorkingHours> . We have worked on this iteration for about 150 hours. Implementing accounts for a rather small percentage of this number. Most of the time we found ourselves thinking and discussing design decisions.

3 Package visible constructors

We decided to make as many constructors as we could package visible only. Doing this denies the creation of the affected objects by any unwanted classes/people efficiently. However, the lack of friend classes and friend packages in Java made it impossible to protect every constructor of every class in such a way.

If such a class or method exists, there will be documentation stating that the method or class in question should never be used outside of the domain layer.

For example: once Machine objects get created, they are added to a MachinePool immediately. The responsibility of assigning every Machine to a MachinePool lays with the MachinePool itself.

4 Creational design patterns

The following problem presented itself to us a couple of times while developing a plausible design for the assignment: what is a *clean* way to allow users to create new objects of certain dynamically changing classes? For example: how can we make sure that the user has the ability to create a new ultra sound scanner without making the domain inconsistent or giving too much responsibility to the user interface?

An answer to this question is implementing a creational design pattern. We decided to use a builder pattern for the creation of new Machines and an abstract factory pattern for the creation of new Users, Treatments, MedicalTests.

The builder mentioned above is nothing like the StandardHospitalBuilder, ExtendedHospitalBuilder, WarehouseBuilder, UserManagerBuilder,... that you may have noticed in the source code of our project. In fact, they're nothing alike. The latter builders are templates for an initialised system to avoid a bunch of code duplication across different testing units.

Without further a due, we will now briefly discuss the creational patterns mentioned in the paragraph above.

4.1 User factories and user manager




Figure 1: The sequence diagram from the add hospital use case demonstrates how user factories are to be used and how they work in the domain layer.

First, let's take a look at the sequence diagram of add new hospital staff above. We can see that the hospital administrator asks a list of all available factories, then selects the kind of factory he would like to create a new staff member from, next sets the required parameters and then lets the system create such a user.

The creation of the actual User object is the responsibility of the UserManager. This makes it easy to make certain that, when the user uses the public controller API, it's impossible to create a new User object that's not kept in a UserManager.

4.2 Machine builder and machine pool

Blood analysers, x-ray scanners and ultra sound scanners all have the same signature in their constructors. They also do not have any specific characteristics other than their types. This is the reason we chose to use a builder design pattern as a solution to the creational problem described in the beginning of this chapter.

It is also important to note that the constructor and build()-method of each builder is package visible only. This means no new MachineBuilder objects can be created without going through the MachinePool. Also the constructors of each of the machine types are package visible only. The combination of these visibility limitations means that, analogous to the case with UserFactory and UserManager, creating a new Machine or MachineBuilder object without registering it in a MachinePool is impossible.

Alas, the visibility in Java is not what it could be. We were forced to make the Machine classes themselves public as we require Machine objects in other packages. In particular, packages that contain TaskDescriptions come to mind, but more on that later.

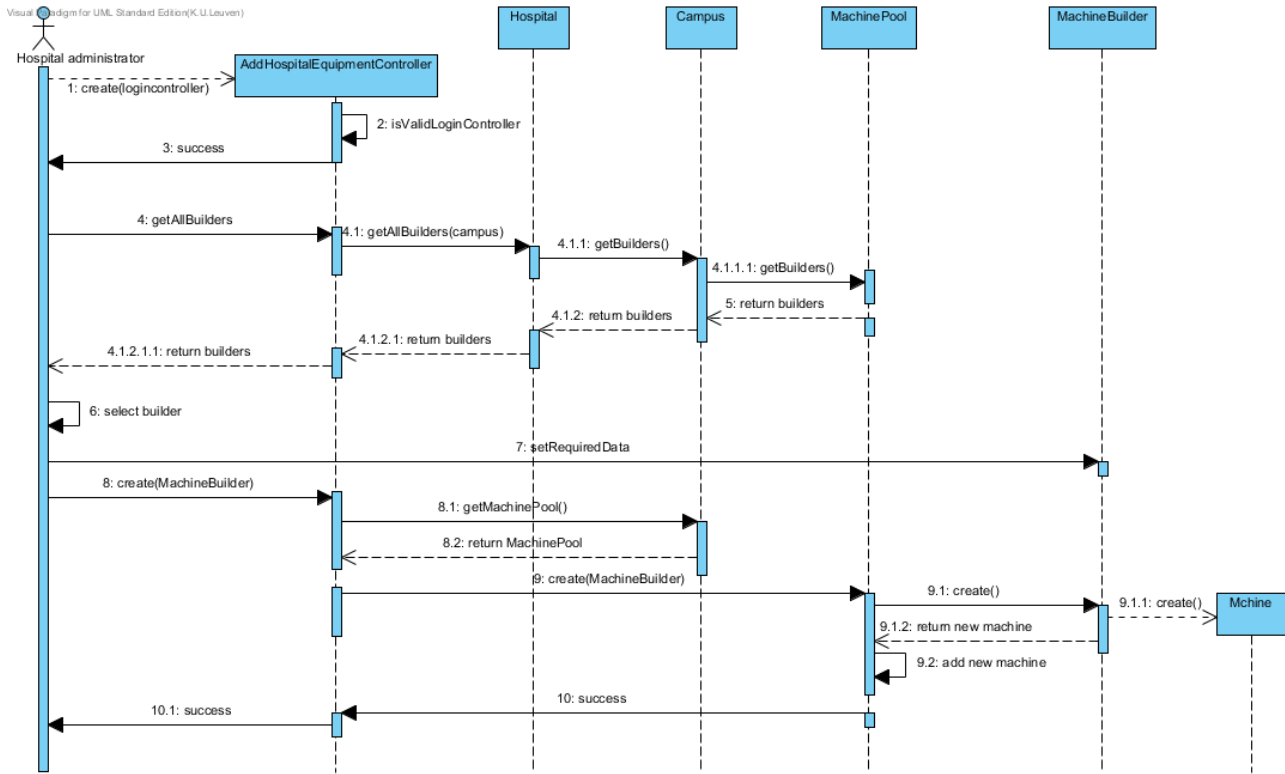


Figure 2: A sequence diagram of the add equipment use case

4.3 Medical test factories and patient files

When a doctor wants to order a new medical test for a patient, he or she will create a new `OrderMedicalTestController`. This controller allows the doctor in question to get all existing medical test factories. These factories are stored in the `Hospital` class and should be added manually at initialisation time before any use cases can be executed.

Once said factories have been acquired from the controller, the doctor can specify the required details in order for the medical test he's about to order to be exactly as he or she would like it to be. The doctor then calls upon the controller one more time, giving it the pre-set `MedicalTestFactory`. The controller passes this factory on to the task manager. The task manager executes the `create()`-method. This method generates a new `MedicalTestDescription`. This description is then stored in a new `Task`. The `Task` is then initialised (which means it will be stored in the patient file of the patient who the medical test is for).

We would have preferred for the `create()`-method to have limited visibility. Once again Java does not provide us with the possibilities to do so. Therefore we have equipped it with the `@Deprecated`-annotation so that, should someone try to call this method, they will see that it's not meant to be called. The documentation of the `create()`-method also informs the reader of this.

Like before, this way of things allows for a reliable domain layer to be constructed and maintained throughout the execution of the program. No `MedicalTest` objects will ever exist without there being a `Task`, as well as a `PatientFile`, to store them. The way `Tasks` are created (by first creating descriptions and then transforming them into `Tasks` in the `TaskManager`) also assures that `Tasks` never exist created without a `TaskManager` keeping track of them. Also, we would like to note that the constructor of `Task` is package visible within the `scheduler.tasks` package to aid in achieving this goal.

4.4 Treatment factories and diagnosis

Treatment factories allow doctors to create treatments that are always linked to diagnose similar to the way medical test factories allow doctors to create medical tests that are always linked to a patient file.

The doctor creates the appropriate controllers and asks them for the treatment factories available in this hospital. Just like is the case with medical test factories, these treatment factories are to be initialised at system initialisation time before any use cases can be ran.

The doctor chooses one of the given treatment factories and sets the required data. Then the doctor gives the factory back to the `PrescribeTreatmentController`, who then in turn, gives it to `TaskManager`. The `TaskManager` then creates a description from the factory and transforms it to a `Task`. The created `Task` gets initialised, which causes it to be linked to the diagnose it was created for.

4.5 Diagnosis and patient files

No `Diagnose` objects should exist without there being a `PatientFile` object to store them. The sequence diagram below demonstrates how assuring this consistency happens when the controllers are used properly:

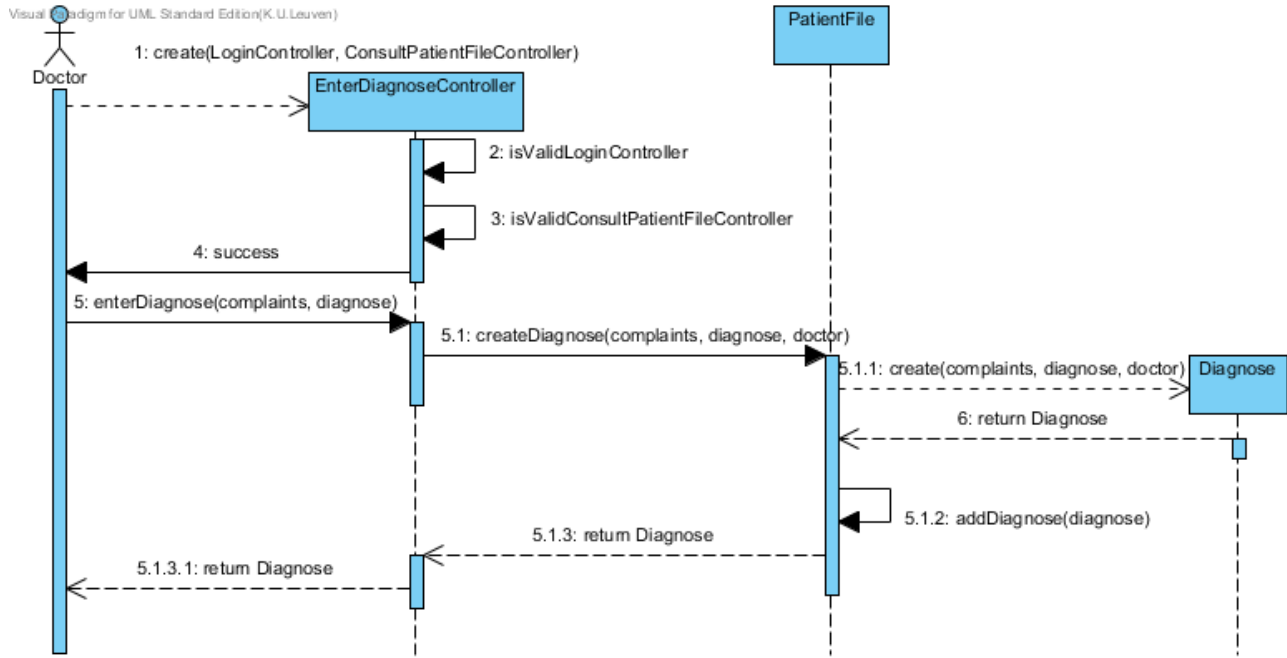


Figure 3: The sequence diagram of the enter diagnose use case demonstrates how no diagnose can be created without a patient file storing it.

4.6 Warehouse builder and campus

Warehouses should not exist without a StockProvider (see later). Also a campus should always have a warehouse. To make this possible we use WarehouseBuilder objects at initialisation time. These warehouse builders are not design pattern builders but merely a way to assure the desired consistency mentioned above at creation time.

5 Task

Task object represent task concepts from the real world. Tasks are created from descriptions and are stored in the TaskManager. Tasks can undergo 2 transformations in their lifetimes. When they are created, they are regarded to as queued tasks. When the opportunity presents itself, the queued task becomes a scheduled task. Once the scheduled task has been executed, it becomes a finished task.

5.1 Generic tasks versus a hierarchy of tasks

Tasks come in many different forms. In the last iteration, we worked with a hierarchy of Task objects. For example: `UnscheduledMedicalTestTask`, `ScheduledMedicalTestTask`, `UnscheduledTreatmentTask`, `ScheduledTreatmentTask`, ... During a meeting with our project advisor, we came to the conclusion that such a hierarchy is not very desirable and violates the GRASP patterns. We then chose to work with a generic Task class: `Task<? extends TaskDescription>`. "Tasks from task descriptions".

The class diagram from Tasks and their descriptions can be found in figure 4.

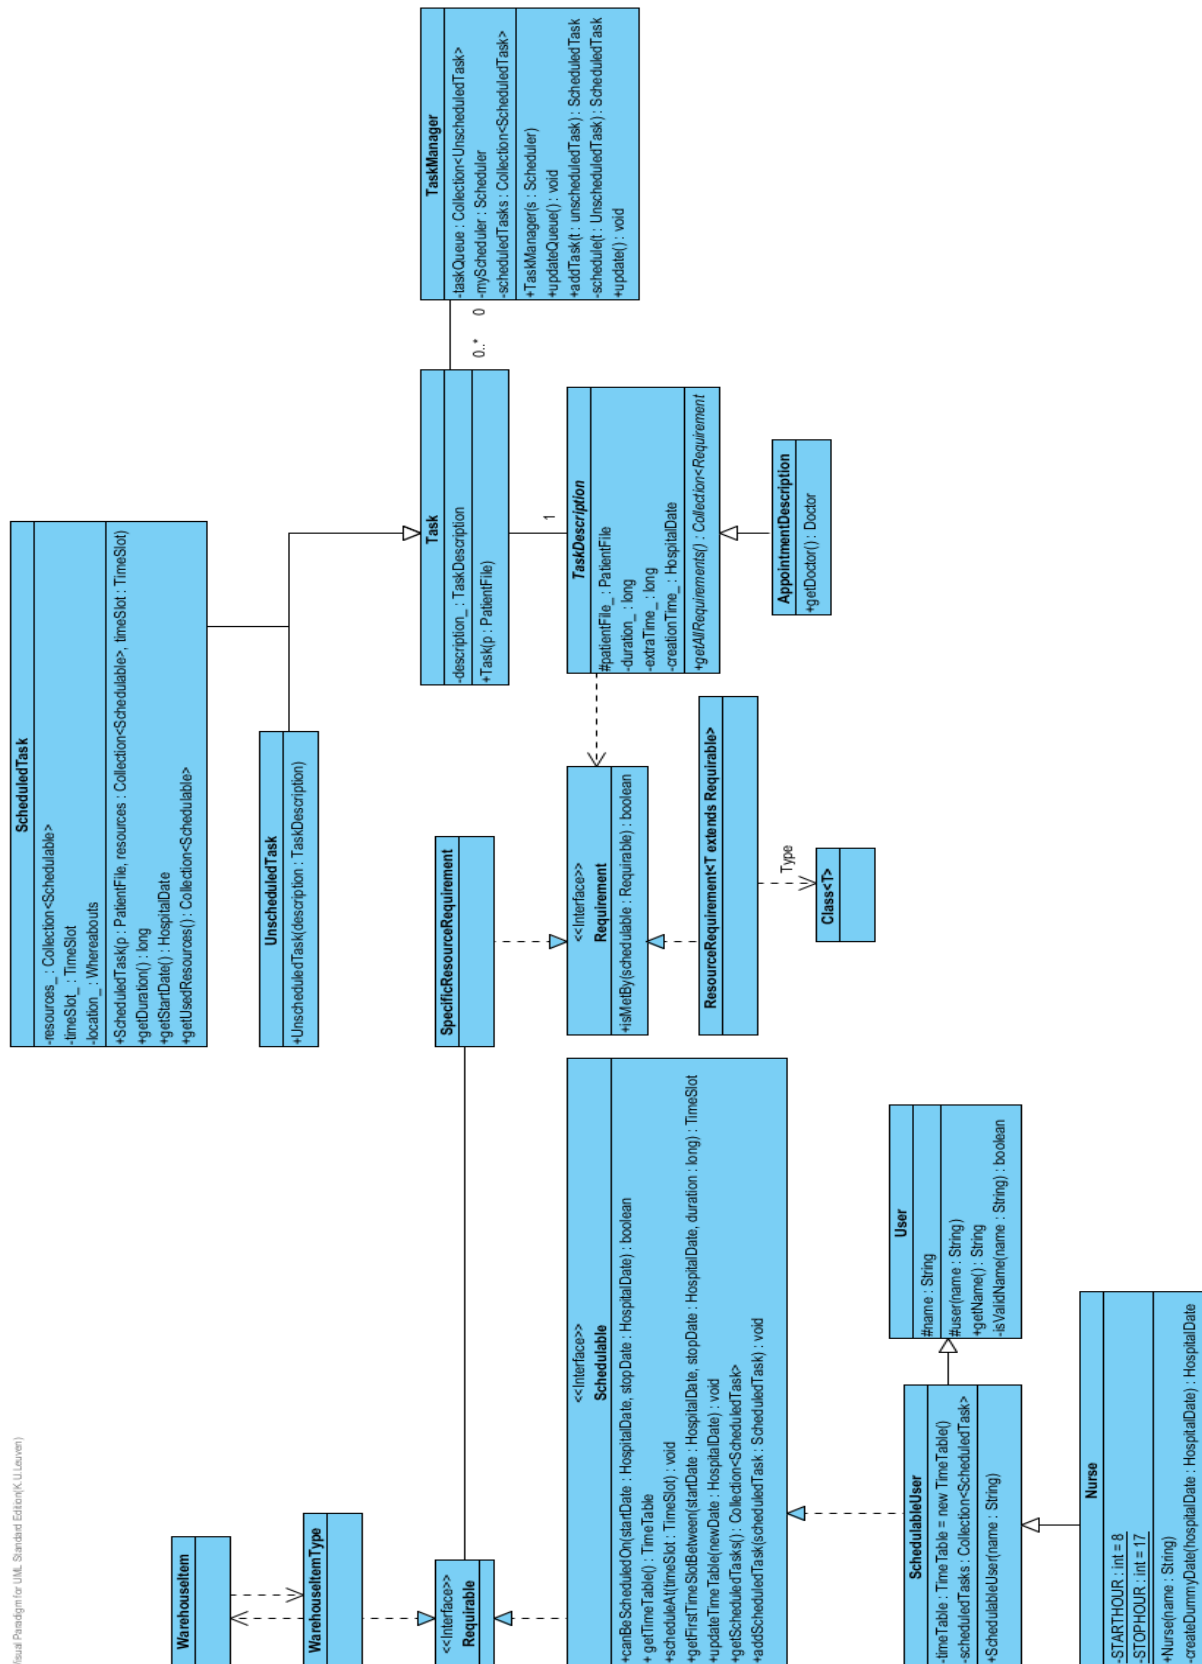


Figure 4: Tasks and their descriptions




Figure 5: Tasks and their states

5.2 The state of a Task

As described in the beginning of this chapter, tasks undergo "transformations". These transformations are implemented with a state pattern.

A task keeps its **TaskState**. The existing **TaskState** objects are: **QueuedState**, **ScheduledState**, **FinishedState**. Each state keeps a **TaskData** object. This object stores information relevant to the current state of the task. For example: a queued task would only keep its description while a scheduled task also keeps all resources involved in the execution of the task, where the task is to take place, etc...

5.3 The description of a Task

We define tasks by their description. Examples of these descriptions are: blood analysis, surgery, medication, appointment,...

These descriptions can be created using the controllers and the factories they can provide the user with. Once created, a description gets stored in the task data of a new queued task. A description contains requirements and conditions specific for that description. These requirements and conditions need to be met in order to schedule the queued task. (see later) You can find a diagram of these states below in figure 5.

6 Warehouse

During this iteration we also made some large refactoring, if not implementing from scratch, from everything related to warehouses. There were a couple of significant problems, like mentioned in the introduction of this report. We managed to successfully find clean solutions for them all.

Let's take a look at the class diagram from warehouse first:

The basic idea is that, using a **WarehouseBuilder**, we connect an **ExpiredItemRemover** and as many **OrderStrategies** to the warehouse as needed.

Visual Paradigm for UML, Standard Edition (K.U. Leuven)

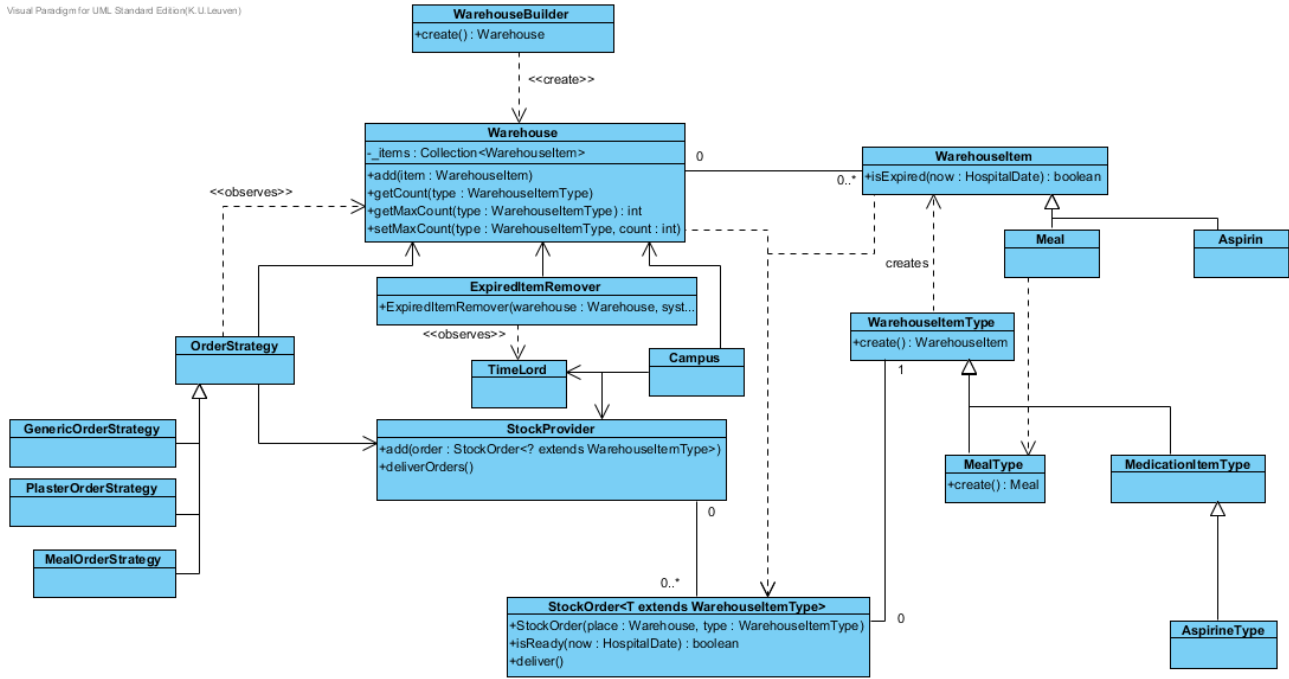


Figure 6: Class diagram from warehouse and its related classes

OrderStrategy is a class that, when the item type it's observing is running low, creates new stock orders for the warehouse automatically.

As different types require different methods of ExpiredItemRemover observes the current system time. Should the time change, the item remover will check if there are items in the warehouse that have expired due to the new system time and will remove them. This removal will, evidently, trigger order strategies to be notified.

All newly created stock orders are kept in the stock provider. The stock provider also observes the current system time. When the time changes, the stock provider will attempt to deliver all deliverable stock orders. Stock orders are deliverable if the required time to pass between order creation and the new system time is sufficient.

The stock provider is created together with the warehouse and both are stored in the campus. Working with only one stock provider per campus makes it easy to create the exact amount of new stock orders required.

The warehouse stores

7 Scheduling

The Scheduling algorithm is built in such a way so that the Scheduler doesn't know anything about the specific classes that are in the the rest of the domain layer. It only knows about abstract schedulables and requirements.

7.1 Description and Requirements

A Description is an abstract class that *describes* a possible task. Such a Description can be an Appointment, a BloodTest, It has all the information it needs to be able to build a task.

This information is contained in the form of requirements. A Requirement is a class such as a Doctor, a WarehouseType, a Nurse, Requirements know what Objects are needed to fulfill this Requirement. There are different Requirements, since not all Requirements are of the same type:

- **SpecificRequirement:** This kind of Requirement contains one Object that needs to be fulfilled. For example a Doctor during an Appointment. This Requirement will only be met by this Doctor.
- **RequirementType:** This kind of Requirement contains an ObjectType. All Objects of this type will fulfill this Requirement. It is possible that there are more than one Object needed to fulfill this Requirement. This information is also stored in this Requirement. An example of this kind of Requirement are Nurses during a xray scan. Since it is not important which Nurse does the scan, only some Nurse object is enough to fulfill this Requirement.
- **DiagnoseCondition:** This kind of Requirement is only fulfilled if the diagnose of the patient is approved.
- **WarehouseItemCondition:** This kind of Requirement can only be fulfilled if a Warehouse has enough items of the kind of type that is stored in this condition. To check whether this condition is met, a HospitalDate and a Location is needed to check this Condition. This Requirement has a collect method, in order to delete all the items it needs from the given Warehouse.

All these Requirements have the same interface, and thus the same methods. One of the methods is *isCrucial()*. A Requirement is crucial if the Requirement needs a Machine or a SchedulableUser. This Requirement can only be met when the admin adds Machines or SchedulableUsers to the Hospital. If this Requirement is not met, the Task can never be scheduled in this Hospital with the current resources. If there are not enough resources available in the hospital to fulfill this Requirement, the Task will never be able to be Scheduled and will be deleted from the system.

7.2 Scheduler and TaskManager

Due to security reasons, it is impossible to create Tasks. A user can only make Descriptions. If a Description is made, the user is obliged to add this Description to the TaskManager. The TaskManager will make a Task out of it and tries to schedule it immediately after that. The scheduling happens with a scheduling algorithm in the Scheduler class. It will try to schedule this task. It is possible that the Scheduler is not able to schedule the Task. In that case, the Scheduler will throw an exception. There are two types of exceptions. Depending on the sort of the exception, the TaskManager will decide whether it will store this Task in a queue or will throw it away.

During this process, the description will be initialised. During this initialisation, the DiagnoseCondition will start observing the Diagnose, in case the state of the Diagnose would ever change. The TaskManager will also have to observe the Task, so the Diagnose will notify the TaskManager indirectly.

If the Scheduler is able to schedule the Task, it will schedule all Schedulables and collect all the WarehouseItems. The state of the Task will be changed.

7.3 Schedulable and Requirable

Schedulables are Objects that can be scheduled. They have a TimeTable.

Requirables are Objects that are needed to forfill a Requirement. This is in fact just an empty interface to generalize all the Objects that are usefull for this.

7.4 HospitalDate

Just like in the previous iteration, we use our own HospitalDate class, since there is currently no good class in java to satisfy our needs

7.5 TimeTable

A TimeTable is a class that represents the individual TimeTable that each Schedulable has. It very much resembles a real life timetable in which one can keep track of ones appointments and todos. A TimeTable consists of a set of TimeSlots. Each TimeSlot in the TimeTable represents an interval in time during which the TimeTable in question is marked as being occupied. A couple of methods in TimeTable that are worth mentioning (as they're not in the class diagram):

- invert() : TimeTable
- getUnion(TimeTable) : TimeTable
- getIntersect(TimeTable) : TimeTable
- eliminateOverlap() : void

Invert() inverts the TimeTable. This means: busy slots become free slots and the other way around. The result of this method is a new TimeTable that is the inverted one of this TimeTable. Evidently the inverted TimeTable of the inverted TimeTable of an original TimeTable, is the original TimeTable again. This method can be very useful to determine all free slots in a TimeTable. getUnion() gets the union of two TimeTables. The union is defined analogous to the union of two collections in mathematics: its a new TimeTable that is marked as busy everywhere where either one or both of the original TimeTables is busy.

getIntersect() gets the intersection of two TimeTables. Like getUnion(), getIntersect() is defined very much like in mathematics: the intersection of two TimeTables is a new TimeTable that is marked as busy where both of the original TimeTables are busy.

eliminateOverlap() is a method that basically merges timeslots that overlap. For example: if a TimeTable contains a slot from 5AM to 8AM and another one from 7AM to 9AM, eliminateOverlap() will result the two original timeslots being reduced to only one: from 5AM to 9AM.

7.6 TimeSlot

A TimeSlot is nothing more than a representation of an interval of time. It has a starting- and stopping point. Its length is in milliseconds.

7.7 TimePoint

As a TimeSlot has a starting and stopping point, TimePoint is the abstract class that represents those points in time. Each TimePoint has a HospitalDate so as to know which point in time it really represents.

It has two child classes: StartTimePoint and StopTimePoint. At first TimePoint held on to an enumeration so that it would know whether it was a starting- or stopping point. Now however, the design is more adaptable: it's easily possible to add a new sort of TimePoint now due to the high cohesion achieved by using two separate classes.

7.8 Impact of having different locations

The different locations had a huge impact on the scheduling algorithm. Instead of just iteration over all the resources, the algorithm had to take into account that not all resources are available everywhere. This has been elegantly solved, by just iterating over all the locations and then trying to schedule over there. The scheduler will choose the best timeslot amongst all the possibilities.

7.9 The location preference for doctors

Because different campuses were introduced, doctors had the option to switch between them. In order to solve this problem, we introduced states. There are two different states. A selected preference state and a back and forth state. The purpose of these two states is handling the TimeTable-calls. This is possible since every state has a LocationTimeTable. All the calls on Doctor involving the gathering of TimeSlots will be overridden and redirected to the states.

7.9.1 getLocationAt()-method

Since every state contains all the information about the location of the Doctor, the state is able to know where the Doctor is at every point in time.

Because we were a little short in time, we didn't have the time to implement these methods.

8 The updated testing scenario

We have changed the testing scenario we handed in on the deadline of the refactoring report drastically. We will now give a short look into the setting of the new scenario and what it does.

8.1 The state of the hospital at initialisation

The hospital has one hospital administrator and two campuses.
Campus 1 contains:

- Three nurses: Jenny, Jenna, Jeffrey;
- Four doctors: Jonathan, Jens, Jelle, Joanne;
- Three x-ray scanners;
- One ultra sound scanner;

- one WarehouseAdministrator: Greg.

Campus 2 contains:

- Two nurses Joy, Janna
- Two doctors Joe, Jennifer
- Three blood analysers;
- One WarehouseAdministrator: Geoff.

The patient files and their history stored in the hospital:

- Stefaan was diagnosed with "cancer" by Jonathan but the diagnose still needs a second opinion from Jennifer. Complaints are "sleepy". Patient has had a blood analysis in the past, created by Jonathan.
- Dieter was diagnosed with "Hypochondria" by doctor Joanne. The diagnose has been approved. Patient's complaints are "Feel sick ALL the time".
- Wouter was just registered in the hospital by nurse Jenna. Complaints are "arm hurts".
- Thibault is currently discharged. Has had the following MedicalTests in the past: one UltraSoundScan, one XRayScan, one BloodAnalysis. Was diagnosed with "Fatal new variant of pneumonia" and had the following treatment carried out successfully: a surgery "Lung transplant".

After this hospital state has been initialised, the test scenario gets executed.

8.2 The events that occur

- Stefaan gets a new MedicalTest: BloodAnalysis. Jennifer disapproves the previous diagnose after the blood analysis has been carried out. The replacement diagnose is "stress". This diagnose gets reviewed and approved by Jonathan. The medication treatment, created by Jennifer, gets scheduled and results in nothing abnormal. Stefaan gets discharged by Jens.
- Dieter gets a brain surgery treatment. This treatment is carried out without problems. After the result of the surgery has been entered, Jens discharges Dieter.
- Wouter gets an appointment with Joe. Joe creates a blood analysis and an x-ray scan for Wouter. After both tests have happened, Joe concludes that Wouter has a fractured arm. The diagnose is approved without needed a second opinion after which a cast gets scheduled. After the result is entered, Wouter gets discharged.
- Thibault enters the hospital at campus 1 and complains: "stomach hurts, jaundice". Nurse Jenna checks him in and creates an appointment with Joe. Joe creates an ultra sound scan and an x-ray scan. From these medical tests, Joe concludes that Thibault's "liver is failing". The treatment is another surgery: "liver transplant". After the treatment has been carried out and the result has been entered, Thibault gets discharged.

- Five new patients come into campus 1 and all of them have the same symptoms "pain in chest and difficulty breathing". After they get their equal appointments, tests, diagnosis ("Heavy metal poisoning") and treatments, two out of five survive while three of them didn't make it.
- Advance time is called a multitude of times in the execution of this scenario.

9 Things that can be done better

9.1 Diagnose state

The lack of the state pattern in Diagnose gave us many problems, we were too late to discover this and thus were unable to implement the pattern here. We think the states UNAPPROVED, APPROVED, DISM would be very helpful in Diagnose objects.

9.2 Comments

We have a lack of commenting on our code, partially due to time pressure and thanks to the many overhauls in our design. The times that when an exception is thrown it is documented thoroughly is rare even in the controllers.

9.2.1 TimeTables & Locations

We noticed that the timetables were very helpful when we were dealing with one campus, but when multiple campi were introduced to the domain they were not sufficient. An ideal solution would probably be a LocationTimeTable that saves the location of a Requirable with the time or a task timetable that keeps track of which task is being executed at what time. This would mean that the interface of Requireable would change and that occupiedOn(HospitalDate):boolean would change to getTaskOn(HospitalDate):Task, all the required info could possibly be derived from that object.

9.3 Factories & consistency

We have a lack of consistency in where the create() method of a factory can and should be used.

9.4 DoctorState

The state pattern in Doctor is used but not yet in such a way that is sufficient, Doctors don't travel between campi. This has high cohesion with the TimeTable location problem that we faced and did not fix.

9.5 Results

We had little problems implementing the results for different task in the hospital, but because we introduced them (as separate objects) too late in our system they barely had any testing.

10 Eclemma and unit tests

We reached a +- 80% code coverage of our domain. We did little border case testing but most of the functionality is implemented. The warehouse package together with the scheduler and task package has the most coverage, we spend a lot of time on testing the timetables since they weren't easy to implement and are one of the most important classes of our system.

11 Capsules

After we've reviewed the design completely, we made all domain layer packages friends with each other. The public API consists of all the methods in the controller-package and the set-methods from the used factories.

The controller package is a friend package of the ui package. The ui package does not have any other friend packages than the controller package.

12 Conclusion

After a lot of hard work with a lot of ups and downs, we feel that we now finally have a system that's designed in a very acceptable way that covers all the bases. Throughout the project several issues concerning team work surfaced and resurfaced. The main one being communication. Explaining what seems evident on the inside your head to other people who don't see things the same way you do can sometimes lead to faulty interpretations or even frustrations.

Nevertheless we learned a lot about object oriented programming, team work, the limitations of Java, design patterns and refactoring.